

Vector databases for embedding systems

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API




Emmanuel Pire

Senior Software Engineer, DataCamp


Limitations of the current approach

- Loading all the embeddings into memory (1536 floats ~ 13kB/embedding)
- Recalculated embeddings for each new query
- Calculating cosine distances for every embedding and sorting is slow and **scales linearly**




Vector databases

- Embedded documents are *stored* and *queried* from the **vector database**




NoSQL Database

- More flexible structure that allows for *faster querying*



SQL/Relational Database


- Structured data into tables, rows, and columns




Components to store

- Embeddings
- Source texts
- Metadata
 - IDs and references
 - Additional data useful for filtering results

Top tip: Don't store the source text as metadata!



The vector database landscape



¹ Image Credit: Yingjun Wu

Which solution is best?

- **Database management:**
 - Managed → more expensive but lowers workload
 - Self-managed → cheaper but requires time and expertise
- **Open source or commercial?**
 - Open source → flexible and cost-effective
 - Commercial → better support, more advanced features, and compliance
- **Data models:** does the type of data lend itself to a particular database type?
- **Specific features:** does your use case depend on specific functionality, such as multi-modal storage?



Let's practice!

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API

Creating vector databases with ChromaDB

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API




Emmanuel Pire

Senior Software Engineer, DataCamp

Installing ChromaDB

- *ChromaDB* is a simple yet powerful vector database
- Two flavors:
 - **Local**: great for development and prototyping
 - **Client/Server**: made for production



Connecting to the database

```
import chromadb

client = chromadb.PersistentClient(path="/path/to/save/to")
```

- Data will be persisted to disk

Creating a collection

- Collections are analogous to tables

```
from chromadb.utils.embedding_functions import OpenAIEmbeddingFunction
```

```
collection = client.create_collection(  
    name="my_collection",  
    embedding_function=OpenAIEmbeddingFunction(  
        model_name="text-embedding-3-small",  
        api_key="..."  
    )  
)
```

- Collections are able to create embeddings automatically

Inspecting collections

```
client.list_collections()
```

```
[Collection(name=my_collection)]
```

Inserting embeddings

Single document

```
collection.add(ids=["my-doc"], documents=["This is the source text"])
```

- IDs must be provided
- Embeddings will be created by the collection!

Multiple documents

```
collection.add(  
  ids=["my-doc-1", "my-doc-2"],  
  documents=["This is document 1", "This is document 2"]  
)
```

Inspecting a collection

Counting documents in a collection

```
collection.count()
```

```
3
```

Inspecting a collection

Peeking at the first 10 items

```
collection.peek()
```

```
{'ids': ['my-doc', 'my-doc-1', 'my-doc-2'],  
 'embeddings': [...], [...], [...]],  
 'documents': ['This is the source text',  
               'This is document 1',  
               'This is document 2'],  
 'metadatas': [None, None, None]}
```

Retrieving items

```
collection.get(ids=["s59"])
```

```
{'ids': ['s59'],  
 'embeddings': None,  
 'metadata': [None],  
 'documents': ['Title: Naruto Shippûden the Movie: The Will of Fire (Movie)\nDescription: When ...'],  
 'uris': None,  
 'data': None}
```

Netflix dataset

Title: Kota Factory (TV Show)

Description: In a city of coaching centers known to train India's finest...

Categories: International TV Shows, Romantic TV Shows, TV Comedies

Title: The Last Letter From Your Lover (Movie)

Description: After finding a trove of love letters from 1965, a reporter sets...

Categories: Dramas, Romantic Movies

Estimating embedding cost

- Embedding model (`text-embedding-3-small`) costs \$0.00002/1k tokens

```
cost = 0.00002 * len(tokens)/1000
```

- Count tokens with the `tiktoken` library
 - `pip install tiktoken`

¹ <https://openai.com/pricing>

Estimating embedding cost

```
import tiktoken

enc = tiktoken.encoding_for_model("text-embedding-3-small")

total_tokens = sum(len(enc.encode(text)) for text in documents)

cost_per_1k_tokens = 0.00002

print('Total tokens:', total_tokens)
print('Cost:', cost_per_1k_tokens * total_tokens/1000)
```

```
Total tokens: 444463
```

```
Cost: 0.00888926
```

Let's practice!

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API

Querying and updating the database

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API



Emmanuel Pire

Senior Software Engineer, DataCamp

Querying the database

"Movies where people sing a lot"

Embed

[0.012688170187175274,
0.013331197202205658,
0.003130610566586256,
0.01171368733048439,
...]

find_n_closest()

Netflix titles

Querying the database

"Movies where people sing a lot"

query()



Netflix titles
collection

The diagram illustrates a database query process. A thick black line starts from the text "Movies where people sing a lot", goes down, then right, and finally down again to form an arrow pointing into a cylinder representing a database. The cylinder is labeled "Netflix titles collection".

Retrieve the collection

```
from chromadb.utils.embedding_functions import OpenAIEmbeddingFunction

collection = client.get_collection(
    name="netflix_titles",
    embedding_function=OpenAIEmbeddingFunction(api_key="...")
)
```

- Must be specify *the same embedding function* used when adding data to the collection

Querying the collection

```
result = collection.query(  
    query_texts=["movies where people sing a lot"],  
    n_results=3  
)  
print(result)
```

```
{  
  'ids': [['s4068', 's293', 's2213']],  
  'embeddings': None,  
  'documents': [['Title: Quién te cantará (Movie)\nDescription: When a near-...',  
    'Title: Quartet (Movie)\nDescription: To save their posh retirement home, ...',  
    'Title: Sing On! Spain (TV Show)\nDescription: In this fast-paced, high-...']],  
  'metadatas': [[None, None, None]],  
  'distances': [[0.350419282913208, 0.36049118638038635, 0.37080681324005127]]  
}
```

`query()` returns a dict with multiple keys:

- `ids` : The ids of the returned items
- `embeddings` : The embeddings of the returned items
- `documents` : The source texts of the returned items
- `metadatas` : The metadatas of the returned items
- `distances` : The distances of the returned items from the query text

```
{'ids': [...],  
 'embeddings': None,  
 'documents': [...],  
 'metadatas': [...],  
 'distances': [...]}
```

```
'ids': [ ['s4068', 's293', 's2213'] ]
```

```
result = collection.query(  
    query_texts=["movies where people sing a lot"],  
    n_results=3  
)
```

- First list corresponds to the first query_text
- Multiple query texts will return multiple lists

```
{ 'ids': [['s4068', 's293', 's2213']],  
  'embeddings': None,  
  'documents': [ ["Title: Quién te cantará (Movie)\nDescription: When a near-drowning leaves a  
famous singer from the '90s with amnesia, she hires a karaoke singer who can imitate her to  
prep her for a comeback tour.\nCategories: Dramas, Independent Movies, International Movies",  
                  'Title: Quartet (Movie)\nDescription: To save their posh retirement home, former opera  
stars plan a gala recital – until the biggest diva among them refuses to sing.\nCategories:  
Comedies, Dramas, Independent Movies',  
                  'Title: Sing On! Spain (TV Show)\nDescription: In this fast-paced, high-energy karaoke  
competition, singers from all walks of life battle it out for up to 30,000 euros!\nCategories:  
International TV Shows, Reality TV, Spanish-Language TV Shows' ]],  
  'metadata': [[None, None, None]],  
  'distances': [[0.350419282913208, 0.36049118638038635, 0.37080681324005127]] }
```

Updating a collection

```
collection.update(  
  ids=["id-1", "id-2"],  
  documents=["New document 1", "New document 2"]  
)
```

- Include *only* the fields to update, other fields will be unchanged
- Collection will automatically create embeddings

Upserting a collection

```
collection.upsert(  
  ids=["id-1", "id-2"],  
  documents=["New document 1", "New document 2"]  
)
```

- If IDs are missing → add them
- If IDs are present → update them

Deleting

Delete items from a collection

```
collection.delete(ids=["id-1", "id-2"])
```

Delete all collections and items

```
client.reset()
```

- **Warning:** this will delete **everything** in the database!

Let's practice!

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API

Multiple queries and filtering

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API



Emmanuel Pire

Senior Software Engineer, DataCamp

Movie recommendations based on multiple datapoints

- Terrifier (id: 's8170')
- Strawberry Shortcake: Berry Bitty Adventures (id: 's8103')

Multiple query texts

```
reference_ids = ['s8170', 's8103']

reference_texts = collection.get(ids=reference_ids)["documents"]

result = collection.query(
    query_texts=reference_texts,
    n_results=3
)
```

Multiple query texts result

```
{'ids': [['s8170', 's6939', 's7000'], ['s8103', 's2968', 's3085']],  
  'embeddings': None,  
  'documents': ['Title: Terrifier (Movie)...',  
                'Title: Hunters: The Art of the Scare (Movie)...',  
                'Title: Horror Story (Movie)...'],  
  ["Title: Strawberry Shortcake: Berry Bitty Adventures (TV Show)...",  
   "Title: Shopkins (TV Show)...",  
   "Title: Rainbow Ruby (TV Show)..."]],  
  'metadatas': [[None, None, None], [None, None, None]],  
  'distances': [[0.00, 0.25, 0.26], [0.00, 0.25, 0.28]]}
```

Adding metadata

```
import csv

ids = []
metadatas = []

with open('netflix_titles.csv') as csvfile:
    reader = csv.DictReader(csvfile)
    for i, row in enumerate(reader):
        ids.append(row['show_id'])
        metadatas.append({
            "type": row['type'],
            "release_year": int(row['release_year'])
        })
```

- Create a list of dicts for the metadatas
- Create a list of IDs to add them to the existing items

Adding and querying metadata

```
collection.update(ids=ids, metadata=metadata)
```

```
result = collection.query(  
    query_texts=reference_texts,  
    n_results=3,  
    where={  
        "type": "Movie"  
    }  
)
```

Where operators

```
where={  
  "type": "Movie"  
}
```

is the same as

```
where={  
  "type": {  
    "$eq": "Movie"  
  }  
}
```

List of operators:

- `$eq` - equal to (string, int, float)
- `$ne` - not equal to (string, int, float)
- `$gt` - greater than (int, float)
- `$gte` - greater than or equal to (int, float)
- `$lt` - less than (int, float)
- `$lte` - less than or equal to (int, float)

Multiple where filters

```
where={
  "$and": [
    {"type":
      {"$eq": "Movie"}
    },
    {"release_year":
      {"$gt": 2020}
    }
  ]
}
```

- `$or` : filter based on *at least* one condition

```
Title: A Classic Horror Story (Movie) [...]  
===  
Title: Nightbooks (Movie) [...]  
===  
Title: Irul (Movie) [...]  
===  
Title: Intrusion (Movie) [...]  
===  
Title: Things Heard & Seen (Movie) [...]  
===  
Title: A StoryBots Space Adventure (Movie) [...]
```

Let's practice!

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API

Congratulations!

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API




Emmanuel Pire

Senior Software Engineer, DataCamp


Chapter 1 - What are Embeddings?

- **Embeddings:** *vector*/numerical representation of text
- Capture the *semantic meaning* of text
- Used OpenAI's Embedding model
- Can use the *cosine distance* to find similar texts
- Unlocks semantic search, recommendation engines, etc.




Chapter 2 - Embeddings for AI Applications


Semantic search and recommendation



Classification



Chapter 3 - Vector Databases



Where next?

Cloud-based, managed vector databases

- **Pinecone:**
 - [Semantic Search with Pinecone](#) (Code-along)
- **Weaviate:**
 - [Vector Databases for Data Science with Weaviate in Python](#) (Code-along)

Frameworks for creating applications

- **LangChain:**
 - [How to Build LLM Applications with LangChain](#) (Tutorial)
 - [Introduction to Large Language Models with GPT & LangChain](#) (Code-along)

Let's practice!

INTRODUCTION TO EMBEDDINGS WITH THE OPENAI API